



JOINs and Set Operators

.NET

A JOIN is the means used to combine columns from one or more tables by using values common to both tables.

[HTTPS://EN.WIKIPEDIA.ORG/WIKI/JOIN_\(SQL\)](https://en.wikipedia.org/wiki/Join_(SQL))

SQL JOIN Statements – Overview


<https://docs.microsoft.com/en-us/sql/relational-databases/performance/joins?view=sql-server-ver15#fundamentals>

JOINS tell SQL how to use data from one table to select rows in a different table.

A *JOIN* statement defines the relationship between the tables by using keywords to:

- Specifying the column from each table to be compared for the *JOIN*.
- Specifying a logical operator (= or <>,) for comparing values from the indicated columns.


Typically, a *JOIN* condition uses a *Foreign Key* from one table and its associated *Primary Key* in the other table.



SQL (INNER) JOIN Statement

<https://docs.microsoft.com/en-us/sql/relational-databases/performance/joins?view=sql-server-ver15#fundamentals>

The **JOIN** (aka **INNER JOIN**) combines with the **WHERE** and **HAVING** search conditions to control which rows are selected from the tables referenced in the **FROM** clause. Specifying the **JOIN** conditions in the **FROM** clause helps separate them from any other search conditions that may be specified in a **WHERE** clause. **INNER JOINS** eliminate the rows that do not match with a row from the other table.



FROM <first_table> <join_type> <second_table> **ON** <join_condition>

This example:

- Specifies 3 unambiguous column names to return from the desired tables,
- Specifies the 2 tables to JOIN,
- Specifies the columns with shared data ON which to JOIN,
- Sets 2 constraints to filter on the results reported


```
SELECT ProductID, Purchasing.Vendor.BusinessEntityID, Name
FROM Purchasing.ProductVendor JOIN Purchasing.Vendor
ON Purchasing.ProductVendor.BusinessEntityID =
    Purchasing.Vendor.BusinessEntityID
WHERE StandardPrice > $10 AND Name LIKE 'F%'
```


SQL FULL JOIN Statement

<https://docs.microsoft.com/en-us/sql/relational-databases/performance/joins?view=sql-server-ver15#fundamentals>
https://www.w3schools.com/sql/sql_join_full.asp

FULL JOINS (aka **OUTER JOIN**) return all rows from at least one of the tables or views mentioned in the **FROM** clause, if those rows meet any **WHERE** or **HAVING** search conditions.

The first table mentioned in the **FULL JOIN / OUTER JOIN** is the "left" table and the second table is the "right" table. When you specify a **LEFT** or **RIGHT OUTER JOIN**, you are referring to the order in which the tables were added to the query and to the order in which they appear in the SQL statement.



This example:

- Specifies 2 unambiguous column names to return from the desired tables,
- Specifies the 2 tables to FULL OUTER JOIN,
- Specifies the columns with shared data ON which to JOIN,
- Sets a constraint to order the result by CustomerName of the Customers table, Ascending (default).


```
SELECT Customers.CustomerName, Orders.OrderID
FROM Customers FULL OUTER JOIN Orders
ON Customers.CustomerID = Orders.CustomerID
ORDER BY Customers.CustomerName;
```

SQL LEFT JOIN Statement

https://www.w3schools.com/sql/sql_join_left.asp

When you specify a **LEFT** or **RIGHT JOIN**, you are referring to the order in which the tables were written in the query. The first table mentioned in the query is the "left" table and the second table is the "right" table.

The **LEFT JOIN** returns all records from the left table (table1), and the matched records from the right table (table2). The result is NULL from the right side if there is no match.



This example:

- Specifies 2 unambiguous column names to return from the desired tables,
- Specifies the 2 tables to LEFT JOIN,
- Specifies the columns with shared data ON which to LEFT JOIN,
- Sets a constraint to order the result by CustomerName of the Customers table, Ascending (default).

```
SELECT Customers.CustomerName, Orders.OrderID
FROM Customers LEFT JOIN Orders
ON Customers.CustomerID = Orders.CustomerID
ORDER BY Customers.CustomerName;
```


SQL RIGHT JOIN Statement

When you specify a **LEFT** or **RIGHT JOIN**, you are referring to the order in which the tables were written in the query. The first table mentioned in the query is the "left" table and the second table is the "right" table.

The **RIGHT JOIN** returns all records from the right table (table2), and the matched records from the left table (table1). The result is NULL from the left side, if there is no match.

This example:

- Specifies 3 unambiguous column names to return from the desired tables,
- Specifies the 2 tables to RIGHT JOIN,
- Specifies the columns with shared data ON which to RIGHT JOIN,
- Sets a constraint to order the result by OrderID of the Orders table, Ascending (default).



```
SELECT Orders.OrderID, Employees.LastName,  
Employees.FirstName  
FROM Orders RIGHT JOIN Employees  
ON Orders.EmployeeID = Employees.EmployeeID  
ORDER BY Orders.OrderID;
```

SQL SELF JOIN


https://www.w3schools.com/sql/sql_join_self.asp

A **SELF JOIN** is exactly like a **JOIN**. The table is joined with a duplicate version of itself, and the result generated.

EX. A table contains all employees. Some are Managers and some are subordinates. Join the subordinate with their manager.

This example:

- Specifies 3 unambiguous column names to return from the desired tables. The AS keyword allows you to designate a unique identifier.
- Specifies the 2 tables to SELF JOIN. No keyword is required, just a comma between the table names.
- Specifies the filter of CustomerID's that are not equal, but the city is the same.
- Sets a constraint to order the result by City name, Ascending (default).



```
SELECT A.CustomerName AS  
CustomerName1, B.CustomerName AS  
CustomerName2, A.City  
  
FROM Customers A, Customers B  
  
WHERE A.CustomerID <> B.CustomerID  
AND A.City = B.City  
  
ORDER BY A.City;
```

<> means != (not equal to)

SQL CROSS JOIN


<https://www.w3resource.com/sql/joins/cross-join.php>

CROSS JOIN produces a *Cartesian Product*. This is a result set which is the number of rows in the first table multiplied by the number of rows in the second table.

If the **WHERE** clause is used with **CROSS JOIN**, it functions like an **INNER JOIN**. You can also use comma-separated column names after **SELECT** and enter comma-separated table names after the **FROM** keyword.

This example:

- Specifies 4 unambiguous column names from the desired tables.
- Specifies the 2 tables FROM which to CROSS JOIN. The first table to multiplied by all items in second table.



```
SELECT foods.item_name,foods.item_unit,  
company.company_name,company.company_city  
FROM foods CROSS JOIN company;
```

SQL Set Operators

<https://learn.microsoft.com/en-us/sql/t-sql/language-elements/operators-transact-sql?view=sql-server-ver16>

Set Operators are used to combine the result sets of two queries. There are three keywords that represent **Set Operators**, **EXCEPT** and **INTERSECT** and **UNION**.

The basic rules for the **Set Operators** are:

- The number and the order of the columns must be the same in all queries.
- The data types must be compatible (implicitly convertible).
- **EXCEPT** and **INTERSECT** return the result sets' column names that are the same as the column names that the query on the operator's left side returns.

```
{ <query_specification> | ( <query_expression> ) }  
{ EXCEPT | INTERSECT }  
{ <query_specification> | ( <query_expression> ) }
```

```
SELECT ProductID  
FROM Production.Product  
INTERSECT  
SELECT ProductID  
FROM Production.WorkOrder ;
```

SQL UNION

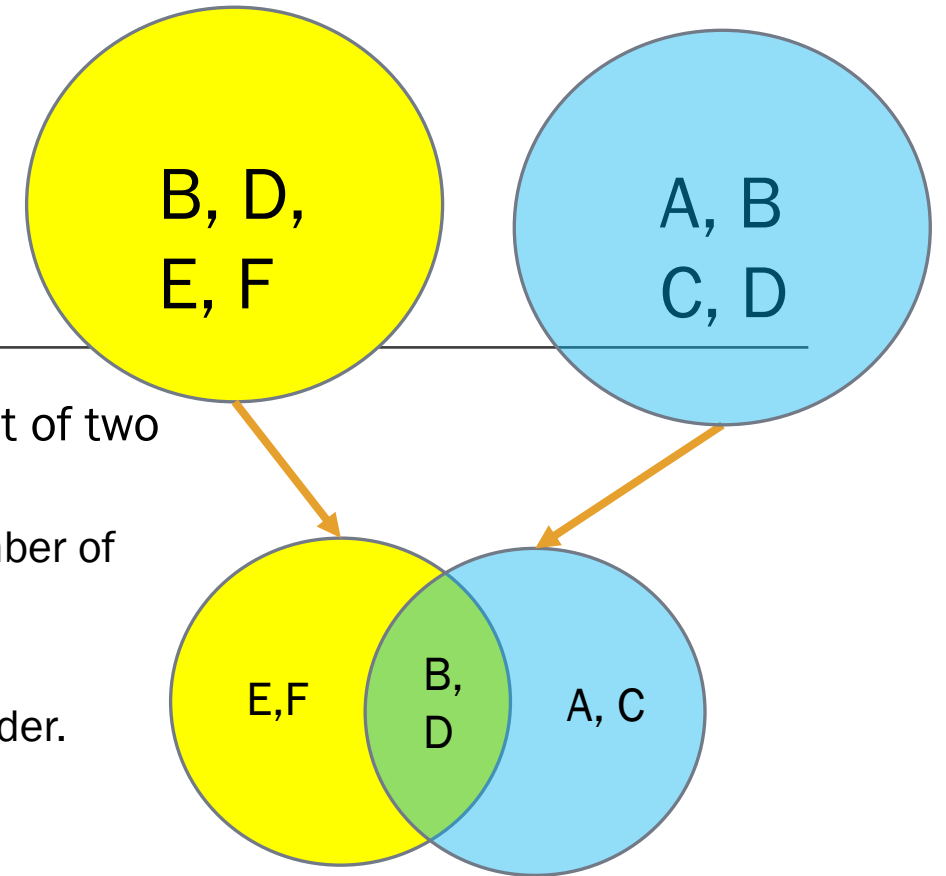
<https://learn.microsoft.com/en-us/sql/t-sql/language-elements/set-operators-union-transact-sql?view=sql-server-ver16>
w3schools.com/sql/sql_union.asp

The **UNION** operator is used to combine the unique result-set of two or more **SELECT** statements.

- Each **SELECT** statement within **UNION** must have the same number of columns.
- The columns must have similar data types.
- The columns in each **SELECT** statement must be in the same order.
- **UNION** selects only distinct values.

This example:

- Creates a complete **SELECT** statement
- Uses the **UNION** keyword
- Creates a separate **SELECT** statement identical to the first but querying a different table
- Sets a constraint to order the result by City name, Ascending (default).



```
SELECT City, Country FROM Customers
WHERE Country='Germany'
UNION
SELECT City, Country FROM Suppliers
WHERE Country='Germany'
ORDER BY City;
```

SQL UNION ALL


<https://learn.microsoft.com/en-us/sql/t-sql/language-elements/set-operators-union-transact-sql?view=sql-server-ver16>
[w3schools.com/sql/sql_union.asp](https://www.w3schools.com/sql/sql_union.asp)

The **UNION ALL** operator is used to combine the full result-set, including duplicates, of two or more **SELECT** statements.

- Each **SELECT** statement within **UNION ALL** must have the same number of columns.
- The columns must have similar data types.
- The columns in each **SELECT** statement must be in the same order.
- **UNION ALL** returns duplicate values.

This example:

- Creates a complete **SELECT** statement
- Uses the **UNION ALL** keyword
- Creates a separate **SELECT** statement identical to the first but querying a different table
- Sets a constraint to order the result by City name, Ascending (default). This result includes duplicates.



```
SELECT City, Country FROM Customers
WHERE Country='Germany'
UNION ALL
SELECT City, Country FROM Suppliers
WHERE Country='Germany'
ORDER BY City;
```

SQL EXCEPT

<https://learn.microsoft.com/en-us/sql/t-sql/language-elements/set-operators-except-and-intersect-transact-sql?view=sql-server-ver16>

EXCEPT:

- Returns any distinct values from the query left of the *EXCEPT* operator (key word). Those values return as long as the right query doesn't return those values as well.

This *EXCEPT* example returns any distinct values from the query left of the *EXCEPT* operator that are NOT also found on the right query. This shows all products that have NOT been ordered.

```
SELECT ProductID FROM Production.Product  
EXCEPT  
SELECT ProductID FROM Production.WorkOrder ;
```


SQL INTERSECT

<https://learn.microsoft.com/en-us/sql/t-sql/language-elements/set-operators-except-and-intersect-transact-sql?view=sql-server-ver16>

INTERSECT:

- Returns any distinct values that are returned by both the query on the left and right sides of the INTERSECT operator.

This *INTERSECT* example returns any distinct values that are returned by both the query on the left (above) and right (below) sides of the *INTERSECT* operator. This shows all products that have been ordered.

```
SELECT ProductID FROM Production.Product  
  
INTERSECT  
  
SELECT ProductID FROM Production.WorkOrder ;
```

JOIN vs UNION

<https://www.geeksforgeeks.org/difference-between-join-and-union-in-sql>

JOIN merges 2 tables horizontally.

UNION merges tables vertically.